



BzGPS

Backblaze GPU Parity Sharder

October 2021

GEFORCE RTX 3080 FAMILY

	GEFORCE RTX 3080 Ti	GEFORCE RTX 3080
NVIDIA CUDA® Cores	10240	8704
Boost Clock (GHz)	1.67	1.71
Base Clock (GHz)	1.37	1.44
Standard Memory Config	12 GB GDDR6X	10 GB GDDR6X
Memory Interface Width	384-bit	320-bit

Fri Oct 8 06:50:30 2021

Processes:							
GPU	GI	CI	PID	Type	Process name	GPU Memory	
	ID	ID				Usage	
No running processes found							

Blocks, threads per block and grids

```
#ifdef USE_GPU
    /// <<< NUMBER_OF_BLOCKS, NUMBER_OF_THREADS_PER_BLOCK>>>
    code_some_shards<<<1, 1>>>(r, shards, r->data_shards, outputs, r->parity_shards, offset, byte_count);
    //code_some_shards2<<<2, 1>>>(r, shards, r->data_shards, outputs, r->parity_shards, offset, byte_count);
    💡 cudaDeviceSynchronize();
#else
    code_some_shards(r, shards, r->data_shards, outputs, r->parity_shards, offset, byte_count);
#endif
]
```

Left (C port/version)

for() loop optimized
away to a CUDA
block

Right (GPU optimized)

```
matrix_t* m = r->parity_rows; // parity rows
for (int byte_index = offset; byte_index < offset + byte_count; byte_index++) {
    for (int output_index = 0; output_index < output_count; output_index++) {
        BYTE* matrixRow = m->data + ((m->columns) * output_index);
        int value = 0;
        for (BYTE input_index = 0; input_index < input_count; input_index++) {
            BYTE a = matrixRow[input_index];
            int addr = (input_index * byte_count) + byte_index;
            BYTE b = input_shards[addr];

#ifdef USE_GPU
            if (a == 0 || b == 0)
            {
                value ^= 0;
            }
            else
            {
                int log_a = LOG_TABLE2[a & 0xFF];
                int log_b = LOG_TABLE2[b & 0xFF];
                int log_result = log_a + log_b;
                value ^= EXP_TABLE2[log_result];
            }
#else
            // the multiply function needs to be configured to be callable from the kernel (GPU
            // as do the log tables above, they must be able to be read from the kernel functio
            // until that's sorted out just deal with it using ifdefs, it's too late tonight to
            value ^= multiply(a, b);
#endif
        }
        int output_addr = (output_index * byte_count) + byte_index;
        outputs[output_addr] = value;
    }
}
```

```
matrix_t* m = r->parity_rows; // parity rows
for (int byte_index = offset; byte_index < offset + byte_count;
    {
        int output_index = blockIdx.x;
        //printf("blockIdx.x: %d blockDim.x: %d threadIdx.x: %d\n",
        BYTE *matrixRow = m->data + ((m->columns) * output_index);
        int value = 0;
        for (BYTE input_index = 0; input_index < input_count; input
            {
                BYTE a = matrixRow[input_index];
                int addr = (input_index * byte_count) + byte_index;
                BYTE b = input_shards[addr];
                if (a == 0 || b == 0)
                {
                    value ^= 0;
                }
                else
                {
                    int log_a = LOG_TABLE2[a & 0xFF];
                    int log_b = LOG_TABLE2[b & 0xFF];
                    int log_result = log_a + log_b;
                    value ^= EXP_TABLE2[log_result];
                }
            }
        int output_addr = (output_index * byte_count) + byte_index;
        outputs[output_addr] = value;
    }
```

Execution Configuration

Let me try to explain this graphically:



In the drawing, each blue rectangle represents a **thread**. Each gray rectangle represents a **block**. The green rectangle represents the **grid**.

Performance Demonstration



Faster?

- Unwind additional loops in the Reed Solomon erasure coding
 - At least the number of parity shards (4x) is possible - may restructure the loops for additional performance beyond 4x
 - Consider using a larger galois field
- Move driver into the kernel to avoid unnecessary copies - think sendfile(); today consider this
 - ReadFile - kernel → userspace
 - WriteFile - userspace → nvidia driver
 - Nvidia driver → GPU card/memory
 - GPU → Nvidia driver
 - ReadFile - nvidia driver → userspace
 - WriteFile - userspace → kernel filesystem write
- See if the file can be memory mapped directly into the GPU; avoiding any reads - or pass in memory as read-only into the GPU avoid synchronization

What do researchers say?

<https://ieeexplore.ieee.org/document/4536322>

Performance results show that the GPU can outperform a modern CPU on this problem by an order of magnitude and also confirm that a GPU can be used to support a system with at least three parity disks with no performance penalty.

-- Accelerating Reed-Solomon coding in RAID systems with GPUs

What's left?

- Implement decoding (file restore) loop
- Turn it into a service
- Consider additional GPU speedups
- Investigate moving it into the kernel

GPUs were meant for more than Bitcoin.

Thank You!!!



Backblaze

